



Assignment Project Exam Help

NoSQL Databases – Part 3

<https://tutorcs.com>

Column-oriented Data Stores

WeChat: cstutorcs

Column-oriented Data Stores¹

- Assignment Project Exam Help**
- Inspired by Google's Bigtable
 - Store **data grouped by columns** (rather than rows) and may have a very large number of columns.



- Other column-oriented data stores
 - Hbase
 - Hypertable

¹ Figure source: S. Harizopoulos, D. Abadi and P. Boncz, Column-Oriented database systems, VLDB 2009



Google's Bigtable - Problem Analysis

- Used by over 60 projects at Google as of 2006, including Web indexing, Google Earth, Google Finance, Orkut, Google Docs, etc.

- **Data types vary** from URLs to web pages to satellite imagery.
- **Latency requirements vary** from backend bulk processing to real-time data processing.
- **Infrastructures vary** from a handful to thousands of servers.

- **Need to scale to a very large size** such as petabytes of data across thousands of commodity servers.

- Most applications **require only single-row transactions.**



Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

Google's Bigtable - Problem Analysis

Assignment Project Exam Help

- Key questions:

- 1 How to represent data? (**expressiveness**)

Key-value pairs are useful but limited

- 2 How to store data? (**scalability**)

Data needs to be distributed across multiple servers

- 3 How to process data? (**efficiency**)

Join on distributed tables needs to be avoided



Google's Bigtable - Problem Analysis

Assignment Project Exam Help

Key questions:

- 1 How to represent data? (**expressiveness**)

Key-value pairs are useful but limited

- 2 How to store data? (**scalability**)

Data needs to be distributed across multiple servers

- 3 How to process data? (**efficiency**)

Join on distributed tables needs to be avoided

- Solution:

WeChat: cstutorcs

One big table

in which **both rows and columns can be split over multiple servers,**
according to **their relatedness.**



Google's Bigtable - Data Structure

Assignment Project Exam Help

- A (big) table is a **multi-dimensional sparse sorted map**.

<https://tutorcs.com>
(row key, column key, timestamp) → value
string string int64 string

- The map is **indexed by** a row key, a column key, and a timestamp.
- Each value in the map is **an uninterpreted array of bytes**.

WeChat: cstutorcs



Google's Bigtable - Data Structure

Assignment Project Exam Help

- A (big) table is a multi-dimensional sparse sorted map
(**row key**, **column key**, **timestamp**) \mapsto **value**

- Example:** a (big) table that stores Web pages

ROW KEY	COLUMN	COLUMN	COLUMN	...
	CONTENTS:	ANCHOR:CNNSI.COM	ANCHOR:MY.LOOK.CA	...
com.cnn.www	$\langle \text{html} \rangle \langle \text{body} \rangle \text{Home} \dots \leftarrow t_1$ 404 Page not found $\leftarrow t_2$ $\langle \text{html} \rangle \langle \text{body} \rangle \text{Inter} \dots \leftarrow t_3$	CNN $\leftarrow t_9$	CNN.com $\leftarrow t_8$...
com.cnn.weather
com.cnn.live
...

- (“com.cnn.www”, “CONTENTS:”, t_1) \mapsto “ $\langle \text{html} \rangle \langle \text{body} \rangle \text{Home} \dots$ ”
- (“com.cnn.www”, “ANCHOR:MY.LOOK.CA”, t_8) \mapsto “CNN.com”



Google's Bigtable - Data Structure (Row Key)

Assignment Project Exam Help

- Row keys are strings of up to 64KB size.
- Row keys are sorted in a lexicographical order.

<https://tutorcs.com>

WeChat: cstutorcs

Row KEY	...
com.cnn.www	...
com.cnn.weather	...
com.cnn.live	...
...	...

- Every read or write of data under a single row key is atomic (regardless of the number of different columns being read or written in the row).



Google's Bigtable - Data Structure (Row Key)

- A table is **dynamically partitioned into tablets** each approximately 100-200 MB in size by default. A tablet can be regarded as a horizontal partition in a table.
- Tablets are **the basic units of distribution and load balancing**, served by **tablet servers**.

	Row KEY	...
tablet ₁	com.cnn.www	...
	com.cnn.weather	...
	com.cnn.live	...

tablet ₂	nz.ac.otago.www	...
	nz.ac.otago.cs	...

...



Google's Bigtable - Data Structure (Row Key)

Assignment Project Exam Help

- Question: Why are the reversed URLs used as row keys?

<https://tutorcs.com>

WeChat: cstutorcs

	Row KEY	...
tablet ₁	com.cnn.www	...
	com.cnn.weather	...
	com.cnn.live	...

tablet ₂	nz.ac.otago.www	...
	nz.ac.otago.cs	...

...



Google's Bigtable - Data Structure (Row Key)

- Applications need to **wisely choose row keys**
 - The ordering of row-keys affects partitioning of rows into tablets.
 - Row ranges with smaller lexicographical distance are split into fewer tablets (good for reads).

<https://tutorcs.com>

WeChat: cstutorcs

	Row Key	
tablet ₁	com.cnn.www	...
	com.cnn.weather	...
	com.cnn.live	...
tablet ₂	nz.ac.otago.www	...
	nz.ac.otago.cs	...

...

- As a result, reads of short row ranges are efficient and typically require communication with only a small number of machines.



Google's Bigtable - Data Structure (Column)

Assignment Project Exam Help

- Columns are **grouped into column families**, i.e., a column family contains columns of related data. A column is named as **family:qualifier**, e.g.,

<https://tutorcs.com>

COLUMN FAMILY 1	COLUMN FAMILY 2		
CONTENTS:	ANCHOR:CNNSI.COM	ANCHOR:MY.LOOK.CA	...
...

WeChat: cstutorcs

- Question:** *Why are columns grouped into column families?*



Google's Bigtable - Data Structure (Column)

Assignment Project Exam Help

- Some properties
 - Column families form **the basic unit of access control**, discerning privileges to read, modify, create column-families, etc.
 - They can be vertically partitioned into different files.
 - Column families need to be defined in the schema (before data can be stored) but **columns within a family can be dynamically changed**.

WeChat: cstutorcs

COLUMN FAMILY 1		COLUMN FAMILY 2	
CONTENTS:	ANCHOR:CNNSI.COM	ANCHOR:MY.LOOK.CA	...
...

- The number of column families should be small (in the hundreds at most).



Google's Bigtable - Data Structure (Timestamp)

Assignment Project Exam Help

- Each cell can contain multiple versions of the same data, indexed by timestamp.

<https://tutorcs.com>

```
<html><body>Home... ←  $t_1$   
404 Page not found ←  $t_2$   
<html><body>Inter... ←  $t_3$ 
```

WeChat: cstutorcs

- Each cell version is a string, e.g., a scalar value.
- Stored in decreasing timestamp order, and thus the most recent version can be **read first**.



Google's Bigtable - Read Operations

Assignment Project Exam Help

```
Scanner scanner(T):
ScanStream *stream;
stream = scanner.FetchColumnFamily("anchor");
stream->SetReturnAllVersions();
scanner.Lookup("com.cnn.www");
for ( ; !stream->Done(); stream->Next() ) {
    printf("%s %s %11d %s\n",
        scanner->RowName(),
        stream->ColumnName(),
        stream->MicroTimestamp(),
        stream->Value());
}
```

<https://tutorcs.com>
WeChat: cstutorcs

Row KEY	COLUMN	COLUMN	COLUMN	...
	CONTENTS:	ANCHOR:CNNSI.COM	ANCHOR:MY.LOOK.CA	...
com.cnn.www	<html><body>Home... ← t_1 404 Page not found ← t_2 <html><body>Inter... ← t_3	CNN ← t_9	CNN.com ← t_8	...
com.cnn.weather
com.cnn.live
...



Google's Bigtable - Write Operations

Assignment Project Exam Help

```
# Open the table
Table *T = OpenOrDie("/bigtable/web/webtable");
```

```
# Write a new anchor and delete an old anchor
```

```
RowMutation r1(T, "com.cnn.www");
r1.Set("anchor:www.c-span.org", "CNN");
r1.Delete("anchor:my.look.ca");
```

```
Operation op;
```

```
Apply(&op, &r1);
```

Row KEY	COLUMN CONTENTS:	COLUMN ANCHOR:CNNSI.COM	COLUMN ANCHOR:MY.LOOK.CA	...
com.cnn.www	$\langle \text{html} \rangle \langle \text{body} \rangle \text{Home} \dots \leftarrow t_1$ 404 Page not found $\leftarrow t_2$ $\langle \text{html} \rangle \langle \text{body} \rangle \text{Inter} \dots \leftarrow t_3$	CNN $\leftarrow t_9$	CNN.com $\leftarrow t_8$...
com.cnn.weather
com.cnn.live
...



Google's Bigtable - Infrastructure Dependencies

- Bigtable is built upon these components:

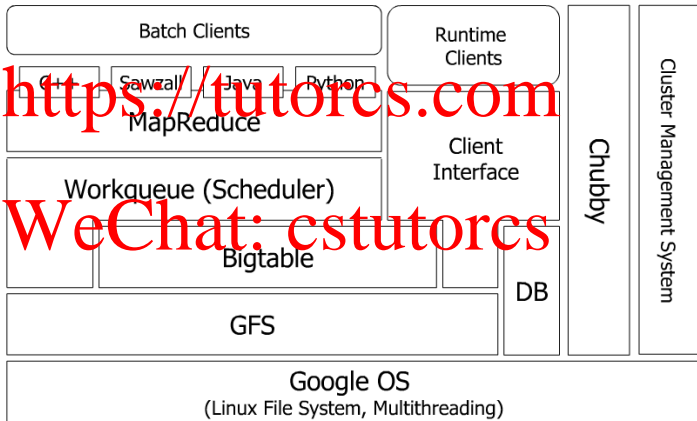
- **Google file system** (GFS): a highly scalable distributed file system
 - e.g., store table data and log.
- **Chubby lock service**: a highly-available and persistent distributed lock service
 - e.g., handles master-election, manage metadata, etc.
- **MapReduce programming model**: a parallel computing model
 - Googles batch processing tool of choice
- **Cluster scheduling system**: a cluster management system
 - e.g., handles failover, monitoring, etc.
- ...

- Similar components are being made available as Open Source by the Apache project Hadoop.



Google's Overall Architecture

- Use **shared-nothing architecture**, consisting of thousands of commodity machines



Google's Bigtable - Summary

Assignment Project Exam Help

- Uses a **shared nothing architecture** to provide scalability over massive data sets:

- **Horizontal partitioning** by range of row keys.
- **Vertical partitioning** by column families

- **Replication**: eventual-consistency replication across datacenters, between multiple BigTable serving setups (master/slave & multi-master)
- Supports **single-row transactions**.
- Supports **only simple queries**.
- Does **not support secondary indices**.