



Assignment Project Exam Help

Query Optimisation

<https://tutorcs.com>

WeChat: cstutorcs



Query Optimisation

Assignment Project Exam Help

- In practice, query optimisers incorporate elements of the following three optimisation approaches:

- **Semantic query optimisation**

Use application specific semantic knowledge to transform a query into the one with a lower cost (they return the same answer).

- **Rule-based query optimisation**

Use heuristic rules to transform a relational algebra expression into an equivalent one with a possibly lower cost.

- **Cost-based query optimisation**

Use a cost model to estimate the costs of plans, and then select the most cost-effective plan.



Semantic Query Optimisation

- Can we use semantic information stored in a database (such as integrity constraints) to optimise queries?

- semantics: “meaning”.

- Recall that, integrity constraints in the relational model include:

- key constraints

- entity integrity constraints

- referential integrity constraints

- domain constraints

- ...

- user-defined integrity constraints

- **Key idea:** Integrity constraints may **not only be utilized to enforce consistency** of a database, but may **also optimise user queries**.



Semantic Query Optimisation

Assignment Project Exam Help

- **Example 1:**

Constraint: The relation *Employee* has the primary key {ssn}.

Query: `SELECT DISTINCT ssn FROM Employee;`

WeChat: cstutorcs

- We can avoid extra costs for duplicate elimination if the existing constraint tells us that tuples in the result will be unique.



Semantic Query Optimisation

Assignment Project Exam Help

● Example 2:

Constraint: No employee can earn more than 200000.

Query:

```
SELECT name  
FROM Employee  
WHERE salary > 300000
```

- We do not need to execute a query if the existing constraint tells us that the result will be empty.



Semantic Query Optimisation

Assignment Project Exam Help

• Example 3:

Constraints: The relation WORKS_ON has the foreign keys:

$[ssn] \subseteq \text{EMPLOYEE}[ssn]$ and $[pno] \subseteq \text{PROJECT}[pnumber]$

Query: `SELECT DISTINCT ssn`

`FROM Works_on INNER JOIN Project`

`on Works_on.pno=Project.pnumber;`

- WeChat: cstutorcs**
- We can reduce the number of joins by executing the following query since both queries always return the same result.

```
SELECT DISTINCT ssn
```

```
FROM Works_on ;
```



Rule-based Query Optimisation

Assignment Project Exam Help

- A rule-based optimisation transforms the RA expression by using a set of heuristic rules that typically improve the execution performance.

- **Key ideas:** apply the most restrictive operation before other operations, which can reduce the size of intermediate results:

<https://tutorcs.com>

- **Push-down selection:**

Apply as early as possible to reduce the number of tuples;

- **Push-down projection:**

Apply as early as possible to reduce the number of attributes.

- **Re-ordering joins:**

Apply restrictive joins first to reduce the size of the result.

- But we must ensure that the resulting query tree gives the same result as the original query tree, i.e., **the equivalence of RA expressions**.

WeChat: cstutorcs



Heuristic Rules

Assignment Project Exam Help

Staff(staffNo, lname, fname, salary, position, branchNo)
Branch(branchNo, name, street, suburb, city)

- There are many heuristic rules for transforming RA expressions, utilized by the query optimiser, such as:

$$(1) \sigma_{\varphi}(\sigma_{\psi}(R)) \equiv \sigma_{\varphi \wedge \psi}(R);$$

$$\sigma_{branchNo='1'}(\sigma_{salary>60000}(Staff)) = \sigma_{branchNo='1' \wedge salary>60000}(Staff)$$

$$(2) \pi_X(\pi_Y(R)) \equiv \pi_X(R) \text{ if } X \subseteq Y;$$

$$\pi_{salary}(\pi_{branchNo, salary}(Staff)) = \pi_{salary}(Staff)$$

$$(3) \sigma_{\varphi}(R_1 \times R_2) \equiv R_1 \bowtie_{\varphi} R_2$$

$$\sigma_{Staff.branchNo=Branch.branchNo}(Staff \times Branch) =$$

$$(Staff) \bowtie_{Staff.branchNo=Branch.branchNo} (Branch)$$



Heuristic Rules

Assignment Project Exam Help

Staff(stid, lname, fname, salary, position, branchNo)
Branch(branchNo, name, street, suburb, city)

<https://tutores.com>

$$(4) \sigma_{\varphi_1}(R_1 \bowtie_{\varphi_2} R_2) \equiv R_2 \bowtie_{\varphi_1 \wedge \varphi_2} R_1$$

$$\sigma_{\text{salary} > 60000}(\text{Staff} \bowtie_{\text{Staff.branchNo} = \text{Branch.branchNo}} (\text{Branch})) =$$

$$(\text{Staff}) \bowtie_{\text{Staff.branchNo} = \text{Branch.branchNo} \wedge \text{salary} > 60000} (\text{Branch})$$

(5) $\sigma_{\varphi}(R_1 \bowtie R_2) \equiv \sigma_{\varphi}(R_1) \bowtie R_2$, if φ contains only attributes in R_1

$$\sigma_{\text{salary} > 60000}(\text{Staff} \bowtie \text{Branch}) \equiv \sigma_{\text{salary} > 60000}(\text{Staff}) \bowtie \text{Branch}$$

(6) $\sigma_{\varphi_1 \wedge \varphi_2}(R_1 \bowtie R_2) \equiv \sigma_{\varphi_1}(R_1) \bowtie \sigma_{\varphi_2}(R_2)$ if φ_1 contains only attributes in R_1 and φ_2 contains only attributes in R_2 .

$$\sigma_{\text{salary} > 60000 \wedge \text{city} = 'Canberra'}(\text{Staff} \bowtie \text{Branch}) =$$

$$(\sigma_{\text{salary} > 60000}(\text{Staff})) \bowtie (\sigma_{\text{city} = 'Canberra'}(\text{Branch}))$$

WeChat: cstutores



Heuristic Rules

Assignment Project Exam Help

Staff(staffNo, lname, fname, salary, position, branchNo)
Branch(branchNo, name, street, suburb, city)

(7) If the join condition involves only attributes in X , we have
 $\pi_X(R_1 \bowtie R_2) \equiv \pi_{X_1}(R_1) \bowtie \pi_{X_2}(R_2)$, where X_i contains attributes in
 both R_1 and R_2 , and ones in both R_i and X , and

$$\pi_{branchNo, position, city}(Staff \bowtie Branch) =$$

$$\pi_{branchNo, position}(Staff) \bowtie (\pi_{branchNo, city}(Branch))$$

(8) If the join condition contains attributes not in X , we have
 $\pi_X(R_1 \bowtie R_2) \equiv \pi_X(\pi_{X_1}(R_1) \bowtie \pi_{X_2}(R_2))$, where X_i contains attributes in
 both in R_1 and R_2 , and ones in both R_i and X

$$\pi_{position, city}(Staff \bowtie Branch) =$$

$$\pi_{position, city}(\pi_{branchNo, position}(Staff) \bowtie (\pi_{branchNo, city}(Branch)))$$



Push-down Selection – Example

Assignment Project Exam Help

- Given the relation schemas:

PERSON(id, first_name, last_name, year_born)

DIRECTOR(id, title, production_year)

MOVIE_AWARD(title, production_year, award_name, year_of_award)

- Query:** List the first and last names of the directors who have directed a movie that has won an 'Oscar' movie award

$\pi_{\text{first_name}, \text{last_name}}(\sigma_{\text{award_name}='Oscar'}((\text{PERSON} \bowtie \text{DIRECTOR}) \bowtie \text{MOVIE_AWARD}))$

- Question:** Can we apply the following rule to optimise the query?

$\sigma_{\varphi}(R_1 \bowtie R_2) \equiv \sigma_{\varphi}(R_1) \bowtie R_2$, if φ contains only attributes in R_1



Push-down Selection – Example

Assignment Project Exam Help

- Given the relation schemas:

PERSON(id, first_name, last_name, year_born)

DIRECTOR(id, title, production_year)

MOVIE_AWARD(title, production_year, award_name, year_of_award)

- Query:** List the first and last names of the directors who have directed a movie that has won an 'Oscar' movie award

$\pi_{first_name, last_name}(\sigma_{award_name='Oscar'}((PERSON \bowtie DIRECTOR) \bowtie MOVIE_AWARD))$

- We would have:

$\pi_{first_name, last_name}((PERSON \bowtie DIRECTOR) \bowtie \sigma_{award_name='Oscar'}(MOVIE_AWARD))$



Push-down Projection – Example

Assignment Project Exam Help

- Given the relation schemas.

PERSON(id, first_name, last_name, year_born)

DIRECTOR(id, title, production_year)

MOVIE_AWARD(title, production_year, award_name, year_of_award)

<https://tutorcs.com>

- Query:** List the first and last names of the directors who have directed a movie that has won an 'Oscar' movie award

$\pi_{first_name, last_name}((PERSON \bowtie DIRECTOR) \bowtie \sigma_{award_name='Oscar'}(MOVIE_AWARD))$

WeChat: cstutores

- Question:** Can we apply the following rule to optimise the query?

$$\pi_X(R_1 \bowtie R_2) \equiv \pi_X(\pi_{X_1}(R_1) \bowtie \pi_{X_2}(R_2)),$$

where X_i contains attributes in both in R_1 and R_2 , and ones in both R_i and X



Push-down Projection – Example

Assignment Project Exam Help

- Given the relation schemas.

PERSON(id, first_name, last_name, year_born)

DIRECTOR(id, title, production_year)

MOVIE_AWARD(title, production_year, award_name, year_of_award)

<https://tutorcs.com>

- Query:** List the first and last names of the directors who have directed a movie that has won an 'Oscar' movie award

$\pi_{first_name, last_name}((PERSON \bowtie DIRECTOR) \bowtie \sigma_{award_name='Oscar'}(MOVIE_AWARD))$

WeChat: cstutores

- we would have:

$\pi_{first_name, last_name}(\pi_{first_name, last_name, title, production_year}(PERSON \bowtie$

$DIRECTOR) \bowtie \pi_{title, production_year}(\sigma_{award_name='Oscar'}(MOVIE_AWARD)))$



A Common Query Pattern (Be Careful)

Assignment Project Exam Help

- A common query pattern is **join-select-project** involving three steps:
 - (1) **join** all the relevant relations,
 - (2) **select** the desired tuples, and
 - (3) **project** on the required attributes.

<https://tutorcs.com>

- This query pattern can be expressed as an RA expression

$$\pi_{A_1, \dots, A_n}(\sigma_{\varphi}(R_1 \times \dots \times R_k)),$$

or as an equivalent SQL statement

```
SELECT DISTINCT  $A_1, \dots, A_n$  FROM  $R_1, \dots, R_k$  WHERE  $\varphi$ ;
```

- Queries falling into this pattern can be **very inefficient**, which may yield huge intermediate result for the joined relations.




A Common Query Pattern (Be Careful)

push-down selection and push-down projection.
Assignment Project Exam Help

$$\pi_{A_1, \dots, A_n}(\sigma_{\varphi}(R_1 \times \dots \times R_k)),$$

<https://tutorcs.com>

WeChat: cstutorcs





Re-ordering Joins - Example

- Given the relation schemas:
`PERSON(id, first_name, last_name, year_born)`

Suppose that it has **10000 tuples**.

`DIRECTOR(id, title, production_year)` with

$[title, production_year] \subseteq MOVIE_AWARD[title, production_year];$
 $[id] \subseteq PERSON[id]$ and

Suppose that it has **100 tuples**.

`MOVIE_AWARD(title, production_year, award_name, year_of_award)`

Suppose that it has **1000 tuples**.

- Example:** Consider the following two RA queries. Which one is better?

- `PERSON ⋈ MOVIE_AWARD ⋈ DIRECTOR`

- `PERSON ⋈ DIRECTOR ⋈ MOVIE_AWARD`



Cost-based Query Optimisation

Assignment Project Exam Help

- A query optimiser does not depend solely on heuristic optimisation. It estimates and compares the costs of different plans.

<https://tutorcs.com>

- It estimates and compares the costs of executing a query using different execution strategies and chooses one with **the lowest cost estimate**.

WeChat: cstutorcs

- The query optimiser needs to **limit the number of execution strategies** to be considered for improving efficiency.



Summary

Assignment Project Exam Help

- In general, there are many ways of executing a query in a database.
- The user expects the result to be returned promptly, i.e., the query should be **processed as fast as possible**.
- But, the burden of optimising queries should not be put on the user's shoulder. **The DBMSs need to do the job!**
- Nonetheless, SQL is not a suitable query language in which queries can be optimised automatically.
- Instead, SQL queries are **transformed into their corresponding RA queries** and optimised subsequently.
- A major advantage of relational algebra is to **make alternative forms of a query easy to explore**.

<https://tutorcs.com>

WeChat: cstutorcs