



Week 10 Workshop - Database Transactions

# Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



## Transactions

- A **transaction** is a sequence of database operations grouped together for execution as a logic unit in a DBMS.

# Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



## Transactions

- A **transaction** is a sequence of database operations grouped together for execution as a logic unit in a DBMS.

Steps	Transaction
	BEGIN TRANSACTION
1	SELECT balance FROM ACCOUNT WHERE name = 'Steve';
2	UPDATE ACCOUNT SET balance = balance-500 WHERE name = 'Steve';
3	SELECT balance FROM ACCOUNT WHERE name = 'Bob';
4	UPDATE ACCOUNT SET balance = balance+500 WHERE name = 'Bob';
5	COMMIT;

<https://tutorcs.com>

WeChat: cstutorcs


## Transactions

- A **transaction** is a sequence of database operations grouped together for execution as a logic unit in a DBMS.

Steps	Transaction
	BEGIN TRANSACTION
1	SELECT balance FROM ACCOUNT WHERE name = 'Steve';
2	UPDATE ACCOUNT SET balance = balance-500 WHERE name = 'Steve';
3	SELECT balance FROM ACCOUNT WHERE name = 'Bob';
4	UPDATE ACCOUNT SET balance = balance+500 WHERE name = 'Bob';
5	COMMIT;

<https://tutors.com>

WeChat: estutors



## Transactions

- What's the difference between database transactions and programs written by a programming language like C, Java, and Python?

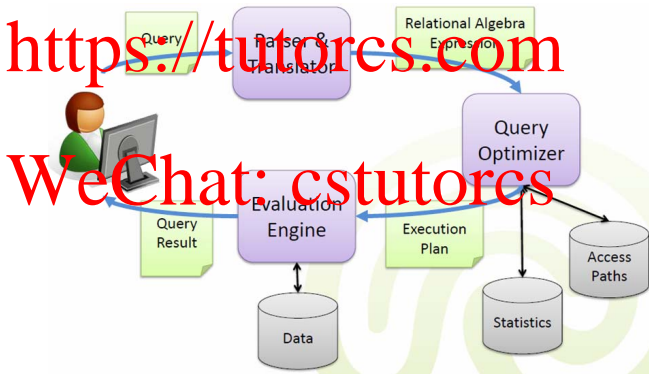
Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

## Transactions

- What's the difference between database transactions and programs written by a programming language like C, Java, and Python?
- How are transactions handled in the query processing?






## Transaction Manager - A Simplified View

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs





## Transactions - ACID Properties

# Assignment Project Exam Help

### Transactions

```
T1 : BEGIN TRANSACTION  
      SELECT ...  
      UPDATE ...  
      COMMIT
```

```
T2 : SELECT ...
```

```
T3 : INSERT ...
```

```
T4 : BEGIN TRANSACTION  
      SELECT ...  
      DELETE ...  
      ABORT
```

<https://tutorcs.com>

WeChat: cstutorcs



## Transactions - ACID Properties

# Assignment Project Exam Help

### Transactions

```
T1 : BEGIN TRANSACTION  
      SELECT ...  
      UPDATE ...  
      COMMIT
```

```
T2 : SELECT ...
```

```
T3 : INSERT ...
```

```
T4 : BEGIN TRANSACTION  
      SELECT ...  
      DELETE ...  
      ABORT
```

### ACID properties

Atomicity

Consistency

Isolation

Durability

<https://tutorcs.com>


WeChat: cstutores

## Transactions - ACID Properties

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs





## Transactions - ACID Properties

# Assignment Project Exam Help

ACID properties

Atomicity

Consistency

Isolation

Durability

<https://tutorcs.com>

WeChat: cstutorcs



## Transactions - ACID Properties

# Assignment Project Exam Help

Transaction Manager

Atomicity

Consistency

Isolation

Durability

<https://tutorcs.com>

WeChat: cstutorcs



## Transactions - ACID Properties

# Assignment Project Exam Help

ACID properties

Atomicity

Consistency

Isolation

Durability

Transaction Manager

Recovery

<https://tutorcs.com>

WeChat: cstutorcs



## Transactions - ACID Properties

# Assignment Project Exam Help

ACID properties

Atomicity

Consistency

Isolation

Durability

Transaction Manager

Recovery


Concurrency  
Control

<https://tutorcs.com>

WeChat: cstutorcs

## Transactions - ACID Properties

Assignment Project Exam Help



Transaction Manager



Consistency is the responsibility of an application developer.

## Transaction Manager - Common Techniques

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs





## Transaction Manager - Common Techniques

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

- **Logging** for recovery – assuring **atomicity/durability** of transactions  
e.g., Write-Ahead Log (WAL) Protocol

## Logging - Introduction

# Assignment Project Exam Help

- A **transaction log** is an append-only file that records changes to objects made by transactions.

<https://tutorcs.com>

WeChat: cstutorcs

## Logging - Introduction

# Assignment Project Exam Help

- A **transaction log** is an append-only file that records changes to objects made by transactions.

<https://tutorcs.com>

- When multiple transactions run concurrently, log records are interleaved.

WeChat: cstutorcs

## Logging - Introduction

# Assignment Project Exam Help

- A **transaction log** is an append-only file that records changes to objects made by transactions.

<https://tutorcs.com>

- When multiple transactions run concurrently, log records are interleaved.

WeChat: cstutorcs

- A transaction log can be implemented as a separate file or set of files in the database.

## Logging - Introduction

# Assignment Project Exam Help

- A **transaction log** is an append-only file that records changes to objects made by transactions.

<https://tutorcs.com>

- When multiple transactions run concurrently, log records are interleaved.

WeChat: cstutorcs

- A transaction log can be implemented as a separate file or set of files in the database.

## Write-Ahead Log (WAL) Protocol

**Assignment Project Exam Help**

Write-Ahead Log (WAL) requires that a record of every change to a database is available while attempting to recover from a crash.

<https://tutorcs.com>

WeChat: cstutorcs



## Write-Ahead Log (WAL) Protocol

# Assignment Project Exam Help

- Write-Ahead Log (WAL) requires that a record of every change to a database is available while attempting to recover from a crash.

• Any change to an object is first recorded in the log, i.e., a record containing both the old and new values for the object.

- A record in the log must be written to persistent storage before committing the transaction.

WeChat: cstutorcs



## Write-Ahead Log (WAL) Protocol

# Assignment Project Exam Help

- Write-Ahead Log (WAL) requires that a record of every change to a database is available while attempting to recover from a crash.

- Any change to an object is first recorded in the log, i.e., a record containing both the old and new values for the object.

- A record in the log must be written to persistent storage before committing the transaction.

WeChat: cstutorcs

- Accordingly, the definition of a **committed transaction** is:

**“A transaction, all of whose log records, including a commit record, have been written to persistent storage”.**





## Write-Ahead Log (WAL) Protocol

- Typical fields in a log record:

LSN  
prevLSN  
transID  
type  
pageID  
length  
offset  
before-image  
after-image

Additional fields for update log records

- Each log record has a unique id called **LSN** (Log Sequence Number).
- **prevLSN** is the LSN of the previous log record written by the same transaction.
- Possible **types** include: update, commit, abort, end, etc.

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



## Write-Ahead Log (WAL) Protocol

- Typical fields in a log record:



- Each log record has a unique id called **LSN** (Log Sequence Number).
- prevLSN** is the LSN of the previous log record written by the same transaction.
- Possible **types** include: update, commit, abort, end, etc.

- Does WAL bring in some benefits for performance?
  - Often results in a significantly reduced number of disk writes
  - Supports one sync against the log file instead of potentially many against the data files
  - Enables online backup and point-in-time recovery

## Transaction Manager - Recovery

# Assignment Project Exam Help

• Key concepts to aid in recovery:

<https://tutorcs.com>

WeChat: cstutorcs

## Transaction Manager - Recovery

# Assignment Project Exam Help

- Key concepts to aid in recovery:

- **Transaction log**: records of database operations

<https://tutorcs.com>

WeChat: cstutorcs

## Transaction Manager - Recovery

# Assignment Project Exam Help

• Key concepts to aid in recovery:

- **Transaction log**: records of database operations

Write-Ahead Log (WAL) – Recovery amounts to either undoing or redoing changes from the log.

WeChat: cstutorcs

## Transaction Manager - Recovery

# Assignment Project Exam Help

• Key concepts to aid in recovery:

- **Transaction log**: records of database operations

Write-Ahead Log (WAL) – Recovery amounts to **either undoing or redoing changes from the log**.

- **Undo** the operations that have not been committed;
- **Redo** the operations that have been committed but not yet been written to disk.

<https://tutorcs.com>  
[WeChat: cstutorcs](https://tutorcs.com)

## Transaction Manager - Recovery

# Assignment Project Exam Help

• Key concepts to aid in recovery:

- **Transaction log**: records of database operations

Write-Ahead Log (WAL) – Recovery amounts to **either undoing or redoing changes from the log**.

- **Undo** the operations that have not been committed;
- **Redo** the operations that have been committed but not yet been written to disk.

- **Checkpoint**: tell the points from which to begin applying transaction logs during database recovery.

(Widely used in practice, but not covered in this course)

## Transaction Manager - Common Techniques

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



- **Logging** for recovery – assuring **atomicity/durability** of transactions  
e.g., Write-Ahead Log (WAL) Protocol
- **Locking** for concurrency control – assuring **isolation** of transactions  
e.g., Two-Phase Locking (2PL) Protocol





## Locking - Introduction

# Assignment Project Exam Help

- A **lock** is associated with an object, e.g., file, table, record, page, etc.
- Two main types of locks:
  - **Read lock** for reading an object by a transaction
  - **Write lock** for writing an object by a transaction

WeChat: cstutorcs



## Locking - Introduction

# Assignment Project Exam Help

- A **lock** is associated with an object, e.g., file, table, record, page, etc.
- Two main types of locks:
  - **Read lock** for reading an object by a transaction
  - **Write lock** for writing an object by a transaction

(Note: there are other types of locks defined by different DBMSs)

WeChat: cstutorcs

## Locking - Introduction

# Assignment Project Exam Help

- A **lock** is associated with an object, e.g., file, table, record, page, etc.
- Two main types of locks:
  - **Read lock** for reading an object by a transaction
  - **Write lock** for writing an object by a transaction

(Note: there are other types of locks defined by different DBMSs)

- Lock compatibility

Lock type	read-lock	write-lock
read-lock	Yes	No
write-lock	No	No



## Two-Phase Locking (2PL) Protocol

# Assignment Project Exam Help

- Locks are handled in two phases:
  - **Expanding**: locks are acquired and no locks are released.
  - **Shrinking**: locks are released and no locks are acquired.

<https://tutorcs.com>


WeChat: cstutorcs

## Two-Phase Locking (2PL) Protocol

# Assignment Project Exam Help

- Locks are handled in two phases:
  - **Expanding**: locks are acquired and no locks are released.
  - **Shrinking**: locks are released and no locks are acquired.

<https://tutorcs.com>



## Two-Phase Locking (2PL) Protocol

**Bad news:**

- 2PL can radically limit interleaving among transactions in some cases ...

<https://tutorcs.com>

WeChat: cstutorcs



## Two-Phase Locking (2PL) Protocol

# Assignment Project Exam Help

Bad news:

- 2PL can radically limit interleaving among transactions in some cases ...
- 2PL may be subject to **deadlocks**, i.e., the mutual blocking of two or more transactions

<https://tutorcs.com>

WeChat: cstutorcs



## Two-Phase Locking (2PL) Protocol

Bad news:

- 2PL can radically limit interleaving among transactions in some cases ...
- 2PL may be subject to **deadlocks**, i.e., the mutual blocking of two or more transactions

Step	$T_1$	$T_2$
1	r-lock(A)	
2	read(A)	
3		r-lock(B)
4		read(B)
5	w-lock(B)	
6	write(B)	
7		w-lock(A)
8		write(A)



## Two-Phase Locking (2PL) Protocol


# Assignment Project Exam Help

Bad news:

- 2PL can radically limit interleaving among transactions in some cases ...
- 2PL may be subject to **deadlocks**, i.e., the mutual blocking of two or more transactions

<https://tutorcs.com>

Step	$T_1$	$T_2$
1	r-lock(A)	
2	read(A)	
3		r-lock(B)
4		read(B)
5	w-lock(B)	
6	write(B)	
7		w-lock(A)
8		write(A)




## Two-Phase Locking (2PL) Protocol

Bad news:

- 2PL can radically limit interleaving among transactions in some cases ...
- 2PL may be subject to **deadlocks**, i.e., the mutual blocking of two or more transactions

<https://tutorcs.com>

Step	$T_1$	$T_2$
1	r-lock(A)	
2	read(A)	
3		r-lock(B)
4		read(B)
5	w-lock(B)	
6	write(B)	
7		w-lock(A)
8		write(A)



- $T_1$  is waiting for  $T_2$  to get a write-lock on  $B$ .  $T_2$  is waiting for  $T_1$  to get a write-lock on  $A$ .

## Two-Phase Locking (2PL) Protocol

# Assignment Project Exam Help

### Good news:

- 2PL makes interleaving safe i.e. guarantee the serializability property for transactions.

<https://tutorcs.com>  
WeChat: cstutorcs

## Two-Phase Locking (2PL) Protocol

# Assignment Project Exam Help

### Good news:

- 2PL makes interleaving safe i.e. guarantee the serializability property for transactions.
- **Serializability** means that a resulting database state is equal to a database state of running transactions serially.

<https://tutores.com>  
WeChat: cstutorcs



## Two-Phase Locking (2PL) Protocol

# Assignment Project Exam Help

### Good news:

- 2PL makes interleaving safe i.e. guarantee the serializability property for transactions.
- **Serializability** means that a resulting database state is equal to a database state of running transactions serially.
- Serializability is the major correctness criterion for concurrent transactions.



## Serializability - Example

- Consider  $A = 200$  and  $B = 500$ , and we have two concurrent transactions:

T1

$A += 100$

$B -= 100$

T2

$A^* = 3$

$B^* = 2$

- Serializable transactions

Assignment Project Exam Help

<https://tutorcs.com>


WeChat: cstutorcs

## Serializability - Example

- Consider  $A = 200$  and  $B = 500$ , and we have two concurrent transactions:



- Serializable transactions







## Serializability - Example

- Consider  $A = 200$  and  $B = 500$ , and we have two concurrent transactions:



- Serializable transactions







## Serializability - Example



- Consider  $A = 200$  and  $B = 500$ , and we have two concurrent transactions:



- Case 1:**  $A=900$  and  $B=800$

- Case 2:**  $A=700$  and  $B=900$

- Serializable transactions



## Serializability - Example


- Consider  $A = 200$  and  $B = 500$ , and we have two concurrent transactions:



- Case 1:**  $A=900$  and  $B=800$

- Case 2:**  $A=700$  and  $B=900$

- Are the following transactions serializable?



## Serializability - Example


- Consider  $A = 200$  and  $B = 500$ , and we have two concurrent transactions:



- Case 1:**  $A=900$  and  $B=800$

- Case 2:**  $A=700$  and  $B=900$

- Are the following transactions serializable?



- Yes.  $A=900$  and  $B=800 \leftrightarrow$  equivalent to Case 1!



## Serializability - Example

- Consider  $A = 200$  and  $B = 500$ , and we have two concurrent transactions:


T1     $A += 100$      $B -= 100$

T2     $A^* = 3$      $B^* = 2$

- Case 1:**  $A=900$  and  $B=800$

- Case 2:**  $A=700$  and  $B=900$

- Are the following transactions serializable?



## Serializability - Example


- Consider  $A = 200$  and  $B = 500$ , and we have two concurrent transactions:



- Case 1:**  $A=900$  and  $B=800$

- Case 2:**  $A=700$  and  $B=900$

- Are the following transactions serializable?



- No.  $A=900$  and  $B=900 \leftrightarrow$  not equivalent to Case 1 or Case 2!



## Problems in Concurrent Transactions

# Assignment Project Exam Help

- If no concurrency control for transactions, some problems may occur:

<https://tutorcs.com>

Lost  
update

write

WeChat: cstutorcs

## Problems in Concurrent Transactions

# Assignment Project Exam Help

- If no concurrency control for transactions, some problems may occur:

<https://tutorcs.com>

**Lost  
update**

write

write

**Dirty  
read**

write

read

**WeChat: cstutorcs**



## Problems in Concurrent Transactions

# Assignment Project Exam Help

- If no concurrency control for transactions, some problems may occur:

<https://tutorcs.com>

**Lost  
update**

write

write

**Dirty  
read**

write

read

**Unrepeatable  
read**

read  
(read)

write

WeChat: cstutorcs





## Problems in Concurrent Transactions

# Assignment Project Exam Help

- If no concurrency control for transactions, some problems may occur:

<https://tutorcs.com>

**Lost  
update**

write  $\rightarrow$  write

**Dirty  
read**

write  $\rightarrow$  read

**Unrepeatable  
read**

read  $\rightarrow$  write  
(read)

**Phantom  
read**

read  $\rightarrow$  write  
(read)

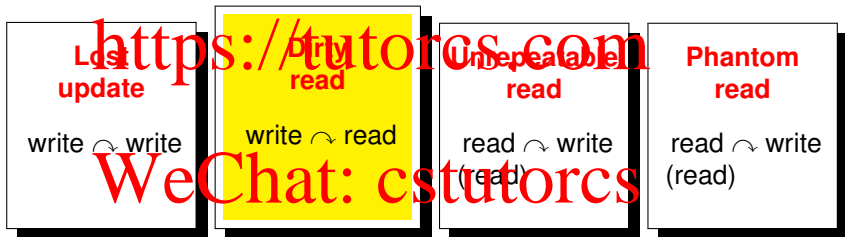
**WeChat: cstutorcs**



## Problems in Concurrent Transactions

# Assignment Project Exam Help

- If no concurrency control for transactions, some problems may occur:




**Dirty write**

write  $\rightarrow$  write

## Problems in Concurrent Transactions

# Assignment Project Exam Help

- If no concurrency control for transactions, some problems may occur:



## The Lost Update Problem - Another Example

- Ben and Amy have the same salary.  $T_1$  sets their salaries to \$80,000, and  $T_2$  sets their salaries to \$90,000.

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



## The Lost Update Problem - Another Example

- Ben and Amy have the same salary.  $T_1$  sets their salaries to \$80,000, and  $T_2$  sets their salaries to \$90,000.

- 1 If executing  $T_1$  and  $T_2$  sequentially,
  - for  $T_1$ ;  $T_2$ , both receive \$90,000.

- for  $T_2$ ;  $T_1$ , both receive \$80,000.

→ Either is acceptable from the transaction viewpoint.

WeChat: cstutorcs



## The Lost Update Problem - Another Example

- Ben and Amy have the same salary.  $T_1$  sets their salaries to \$80,000, and  $T_2$  sets their salaries to \$90,000.

- If executing  $T_1$  and  $T_2$  sequentially,
  - for  $T_1$ ;  $T_2$ , both receive \$90,000.
  - for  $T_2$ ;  $T_1$ , both receive \$80,000.

→ Either is acceptable from the transaction viewpoint.

- If executing  $T_1$  and  $T_2$  concurrently, we may have:

	$T_1$	$T_2$
1	write(A) (A:=80000)	
2		write(A) (A:=90000)
3		write(B) (B:=90000)
4		commit
5	write(B) (B:=80000)	
6	commit	

→ It is not acceptable!



## The Dirty Read Problem - Another Example

- Both Ben and Amy are rewarded a bonus \$5,000 and a pay rise 5%.  $T_1$  increases her salary with \$5,000 and  $T_2$  increments their salaries by 5%.

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



## The Dirty Read Problem - Another Example

- Both Ben and Amy are rewarded a bonus \$5,000 and a pay rise 5%.  $T_1$  increases their salaries with \$5,000 and  $T_2$  increments their salaries by 5%.
- 1 If executing  $T_1$  and  $T_2$  sequentially, they would have .... Also,  $T_1$  or  $T_2$  could abort for some reasons.  $\leftrightarrow$  all are acceptable from the transaction viewpoint.

<https://tutorcs.com>

WeChat: cstutorcs



## The Dirty Read Problem - Another Example

- Both Ben and Amy are rewarded a bonus \$5,000 and a pay rise 5%.  $T_1$  increases her salaries with \$5,000 and  $T_2$  increments their salaries by 5%.

- If executing  $T_1$  and  $T_2$  sequentially, they would have .... Also,  $T_1$  or  $T_2$  could abort for some reasons.  $\hookrightarrow$  all are acceptable from the transaction viewpoint.

- If executing  $T_1$  and  $T_2$  concurrently, we may have:

	$T_1$	$T_2$
1	read(A)	
2	write(A) ( $A := A + 5000$ )	
3		read(A)
4	read(B)	
5	write(B) ( $B := B + 5000$ )	
6	abort	
7		write(A) ( $A := A + A \times 5\%$ )
8		read(B)
9		write(B) ( $B := B + B \times 5\%$ )
10		commit

$\hookrightarrow$  It is not acceptable!



## The Dirty Write Problem - An Example

- Both Ben and Amy are rewarded a bonus \$5,000 and a pay rise 5%.  $T_1$  increases her salaries with \$5,000 and  $T_2$  increments their salaries by 5%.

- If executing  $T_1$  and  $T_2$  sequentially, they would have .... Also,  $T_1$  or  $T_2$  could abort for some reasons.  $\hookrightarrow$  all are acceptable from the transaction viewpoint.

- If executing  $T_1$  and  $T_2$  concurrently, we may have:

	$T_1$	$T_2$
1	read(A)	
2	write(A) ( $A := A + 5000$ )	
3		read(A)
4		write(A) ( $A := A \times 5\%$ )
5	read(B)	
6	write(B) ( $B := B + 5000$ )	
7	abort	
8		read(B)
9		write(B) ( $B := B \times 5\%$ )
10		commit

$\hookrightarrow$  It is not acceptable!

## The Unrepeatable Read Problem - Another Example

# Assignment Project Exam Help

- Amy and Ben are using a website to book flight tickets to Brisbane.

<https://tutorcs.com>

WeChat: cstutorcs

## The Unrepeatable Read Problem - Another Example

# Assignment Project Exam Help

- Amy and Ben are using a website to book flight tickets to Brisbane.
  - Amy signs on first to see that only one ticket is left, and finds it expensive.

<https://tutorcs.com>

WeChat: cstutorcs



## The Unrepeatable Read Problem - Another Example

# Assignment Project Exam Help

- Amy and Ben are using a website to book flight tickets to Brisbane.
  - Amy signs on first to see that only one ticket is left, and finds it expensive.
  - Amy takes time to decide. Ben signs on later and also finds one ticket left, orders it instantly, and logs off.

WeChat: cstutorcs



## The Unrepeatable Read Problem - Another Example

# Assignment Project Exam Help

- Amy and Ben are using a website to book flight tickets to Brisbane.
  - Amy signs on first to see that only one ticket is left, and finds it expensive.
  - Amy takes time to decide. Ben signs on later and also finds one ticket left, orders it instantly, and logs off.
  - Amy decides to buy a ticket, and finds no tickets left.

WeChat: cstutorcs



## The Unrepeatable Read Problem - Another Example

# Assignment Project Exam Help

- Amy and Ben are using a website to book flight tickets to Brisbane.
  - Amy signs on first to see that only one ticket is left, and finds it expensive.
  - Amy takes time to decide. Ben signs on later and also finds one ticket left, orders it instantly, and logs off.
  - Amy decides to buy a ticket, and finds no tickets left.

<https://tutorcs.com>

WeChat: cstutorcs

	$T_1$ (from Amy)	$T_2$ (from Ben)
1	read(X)	
2		read(X)
3		write(X) ( $X := X - 1$ )
4		commit
5	read(X)	



## The Unrepeatable Read Problem - Another Example

# Assignment Project Exam Help

- Amy and Ben are using a website to book flight tickets to Brisbane.
  - Amy signs on first to see that only one ticket is left, and finds it expensive.
  - Amy takes time to decide. Ben signs on later and also finds one ticket left, orders it instantly, and logs off.
  - Amy decides to buy a ticket, and finds no tickets left.

	$T_1$ (from Amy)	$T_2$ (from Ben)
1	read(X)	
2		read(X)
3		write(X) ( $X := X - 1$ )
4		commit
5	read(X)	

- This situation can never arise in a serial execution of  $T_1$  and  $T_2$ .



## The Phantom Read Problem - Another Example

- Amy is 30 years old, but her age in the table players is mistakenly recorded as 40. Ben is 28 years old and his age is correctly recorded in players.

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs



## The Phantom Read Problem - Another Example

- Amy is 30 years old, but her age in the table players is mistakenly recorded as 40. Ben is 28 years old and his age is correctly recorded in players.
- Two concurrent transactions:

```
T1: SELECT * FROM players  
      WHERE age<32;  
      ...  
      SELECT * FROM players  
      WHERE age<32;  
      COMMIT;
```

```
T2: UPDATE players  
      SET age=30  
      WHERE rating=8 and name='Amy';  
      COMMIT;
```

WeChat: cstutorcs



## The Phantom Read Problem - Another Example

- Amy is 30 years old, but her age in the table players is mistakenly recorded as 40. Ben is 28 years old and his age is correctly recorded in players.
- Two concurrent transactions:

```
T1: SELECT * FROM players
      WHERE age<32;
      ...
      SELECT * FROM players
      WHERE age<32;
      COMMIT;
```

```
T2: UPDATE players
      SET age=30
      WHERE rating=8 and name='Amy';
      COMMIT;
```

WeChat: cstutorcs

	T <sub>1</sub>	T <sub>2</sub>
1	read(players)	
2		read(players)
3		write(players)
4		commit
5	read(players)	
6	commit	



## The Phantom Read Problem - Another Example

- Amy is 30 years old, but her age in the table players is mistakenly recorded as 40. Ben is 28 years old and his age is correctly recorded in players.
- Two concurrent transactions:

```
T1: SELECT * FROM players
    WHERE age<32;
    ...
    SELECT * FROM players
    WHERE age<32;
    COMMIT;
```

```
T2: UPDATE players
    SET age=30
    WHERE rating=8 and name='Amy';
    COMMIT;
```

WeChat: cstutorcs

	$T_1$	$T_2$
1	read(players)	
2		read(players)
3		write(players)
4		commit
5	read(players)	
6	commit	

- This situation can never arise in a serial execution of  $T_1$  and  $T_2$ .

## Discussion

# Assignment Project Exam Help

What are the differences between 'unrepeatable read' and 'phantom read'?

<https://tutorcs.com>

WeChat: cstutorcs



## Discussion

# Assignment Project Exam Help

What are the differences between 'unrepeatable read' and 'phantom read'?

- **Unrepeatable read**

- Executing the same SELECT twice yields the same tuples, but attribute values might be different;
- May occur when reading objects that are affected by **UPDATE** from another transaction;
- Can be prevented using record-level locking.

<https://tutorcs.com>  
WeChat: cstutorcs



## Discussion

# Assignment Project Exam Help

What are the differences between 'unrepeatable read' and 'phantom read'?

- **Unrepeatable read**

- Executing the same SELECT twice yields the same tuples, but attribute values might be different;
- May occur when reading objects that are affected by **UPDATE** from another transaction;
- Can be prevented using record-level locking.

- **Phantom read**

- Executing the same SELECT twice yields two different sets of tuples;
- May occur when querying a set of tuples that are affected by **INSERT/DELETE/UPDATE** from another transaction;
- Can be prevented using table-level locking.



## What Should We lock?

- Consider the following two concurrent transactions again:

```
T1: SELECT * FROM players
    WHERE age < 32;
...
SELECT * FROM players
    WHERE age < 32;
COMMIT;
```

```
T2: UPDATE players
    SET age = 30
    WHERE rating = 8 and name = 'Amy';
COMMIT;
```

- What objects should the DBMS lock in order to avoid the phantom read problem?

- Table-level locks**

- e.g., read-lock on players for  $T_1$ , write-lock on players for  $T_2$

- Record-level locks**

- e.g., read-lock on every record with age < 32 for  $T_1$ , write-lock on every record with rating = 8 and name = 'Amy' for  $T_2$

- ...

<https://tutorcs.com>

WeChat: cstutorcs





## Transaction Support in SQL

Assignment Project Exam Help

- An explicit transaction may have no **BEGIN TRANSACTION** statement but must be ended with either **COMMIT** or **ABORT (ROLLBACK)** statement.
- When no explicit transaction statements are given, each single SQL statement is considered to be a transaction.

<https://tutorcs.com>

WeChat: cstutorcs



## Transaction Support in SQL

Assignment Project Exam Help

- An explicit transaction may have no **BEGIN TRANSACTION** statement but must be ended with either **COMMIT** or **ABORT (ROLLBACK)** statement.
- When no explicit transaction statements are given, each single SQL statement is considered to be a transaction.
- To give programmers more control over transaction overhead, SQL allows them to specify **isolation level**, i.e., the degree of interference that a transaction is prepared to tolerate on concurrent transactions.

<https://tutorcs.com>  
WeChat: cstutorcs



## Transaction Support in SQL

Assignment Project Exam Help

- An explicit transaction may have no **BEGIN TRANSACTION** statement but must be ended with either **COMMIT** or **ABORT (ROLLBACK)** statement.
- When no explicit transaction statements are given, each single SQL statement is considered to be a transaction.
- To give programmers more control over transaction overhead, SQL allows them to specify **isolation level**, i.e., the degree of interference that a transaction is prepared to tolerate on concurrent transactions.

Key idea:

WeChat: cstutorcs

To trade off **consistency** (i.e., increased risk of violating database integrity) with **performance** (i.e., greater concurrent access to data)



## Isolation Levels

# Assignment Project Exam Help

- SQL-92 defines four isolation levels:

1 Read Uncommitted

2 Read Committed

3 Repeatable Reads

4 Serializable

- To specify an isolation level, e.g.,

```
SET TRANSACTION ISOLATION LEVEL serializable;
```

- The SQL standard does not impose a specific locking scheme or mandate particular behaviors.

## Isolation Levels

# Assignment Project Exam Help

The intention is to prohibit certain problems:

Isolation Level	Dirty Read	Unrepeatable Read	Phantom Read
READ UNCOMMITTED	Yes	Yes	Yes
READ COMMITTED	No	Yes	Yes
REPEATABLE READ	No	No	Yes
SERIALIZABLE	No	No	No

WeChat: cstutorcs

<sup>1</sup> [https://drtom.ch/posts/2011/11/12/The\\_Lost\\_Update\\_Problem\\_-\\_Part\\_1/](https://drtom.ch/posts/2011/11/12/The_Lost_Update_Problem_-_Part_1/)



## Isolation Levels

# Assignment Project Exam Help

- The intention is to prohibit certain problems:

Isolation Level	Dirty Read	Unrepeatable Read	Phantom Read
READ UNCOMMITTED	Yes	Yes	Yes
READ COMMITTED	No	Yes	Yes
REPEATABLE READ	No	No	Yes
SERIALIZABLE	No	No	No


- Different DBMSs implement isolation levels differently.
- The isolation level required for **Lost Update** is debatable (depending on a DBMS's implementations). But in general, it may require the highest level **SERIALIZABLE** to prevent it.<sup>1</sup>

<sup>1</sup> [https://drtom.ch/posts/2011/11/12/The\\_Lost\\_Update\\_Problem\\_-\\_Part\\_1/](https://drtom.ch/posts/2011/11/12/The_Lost_Update_Problem_-_Part_1/)

## Isolation Levels - Concurrency Control

Assignment Project Exam Help

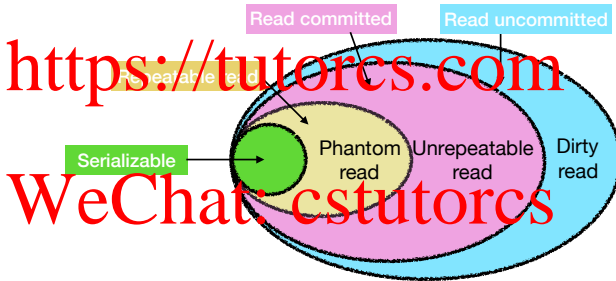
- A DBMS provides different levels of isolation → different degrees of concurrency control to prevent different problems.



## Isolation Levels - Concurrency Control

Assignment Project Exam Help

• A DBMS provides different levels of isolation → different degrees of concurrency control to prevent different problems.



- Concurrency control is **NOT** binary in a database system.







## Isolation Levels - Read Uncommitted

- **Read Uncommitted** is the least restrictive isolation level.
- One transaction can see changes made by other transactions which are not yet committed. This can be quite dangerous.
- Use it when executing queries over read-only data or if it does not matter whether a query returns uncommitted data.





## Isolation Levels - Read Committed

- **Read Committed**: One transaction only sees committed changes by other transactions.
- It is the most commonly used isolation level in database applications.
- Use it when you want to maximize concurrency between applications but do not want queries to see uncommitted data.





## Isolation Levels - Repeatable Reads

# Assignment Project Exam Help

- **Repeatable Reads:** The objects touched by a transaction are locked and cannot be updated or deleted by a concurrent transaction.
- Use it when you want some level of concurrency between applications but do not expect individual objects to be changed during a transaction.


<https://tutorcs.com>  
WeChat: cstutorcs





## Isolation Levels - Serializable

- Serializable is the highest isolation level. All transactions are totally isolated from other transactions. It is safe but may cause significant performance hit.
- Use it when you want some level of concurrency between applications but do not expect that a query returns different sets of results when running at different times.





## Locks Taken by SQL Server for Isolation Levels <sup>2</sup>

# Assignment Project Exam Help



<https://tutorcs.com>

WeChat: [tutorcs](#)

<sup>2</sup><http://michaeljswart.com/2012/06/visualizing-transaction-isolations-for-sql-server/>

## Wrap-up - Isolation Levels

# Assignment Project Exam Help

- A lower isolation level increases the ability of many users to access data **at the same time**, but also increases the number of concurrency effects (such as dirty reads or lost updates) users might encounter.
- Conversely, a higher isolation level **reduces the types of concurrency effects that users may encounter**, but requires more system resources and increases the chances that one transaction will block another.
- Choosing the appropriate isolation level depends on **balancing**  
**WeChat: cstutorcs**
  - **the data integrity requirements of the application**  
against
  - **the overhead of each isolation level.**

# Assignment Project Exam Help

<https://tutorcs.com>  
**Research Topics**

WeChat: cstutorcs

## Research Topics

# Assignment Project Exam Help

- This is an active research area covering many interesting research topics.
- Historically, much of the work has been done in the context of relational database systems.
- However, the ideas in general are independent of whether the underlying system is a relational database system or something else.

<https://tutorcs.com>

WeChat: cstutorcs

- Distributed database systems
- Graph database systems
- Document-oriented database systems
- ...



## Research Topics

- Distributed transactions

